

Import necessary libraries

```
import tensorflow as tf
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

Load Training Data

```
training_set = tf.keras.utils.image_dataset_from_directory(
    'train',
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=32,
    image_size=(128, 128),
    shuffle=True,
    seed=None,
    validation_split=None,
    subset=None,
    interpolation="bilinear",
    follow_links=False,
    crop_to_aspect_ratio=False
)
```

Found 70295 files belonging to 38 classes.

Load Validation Data

```
validation_set = tf.keras.utils.image_dataset_from_directory(
    'valid',
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=32,
    image_size=(128, 128),
    shuffle=True,
    seed=None,
    validation_split=None,
    subset=None,
    interpolation="bilinear",
    follow_links=False,
    crop_to_aspect_ratio=False
)
```

Found 17572 files belonging to 38 classes.

Initialize the Sequential model

```
cnn = tf.keras.models.Sequential()
```

Add convolutional and pooling layers to the model

```
cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, padding='same',
    , activation='relu', input_shape=[128, 128, 3]))
cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

cnn.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, padding='same',
    , activation='relu'))
cnn.add(tf.keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

cnn.add(tf.keras.layers.Conv2D(filters=128, kernel_size=3, padding='same',
    , activation='relu'))
cnn.add(tf.keras.layers.Conv2D(filters=128, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

cnn.add(tf.keras.layers.Conv2D(filters=256, kernel_size=3, padding='same',
    , activation='relu'))
cnn.add(tf.keras.layers.Conv2D(filters=256, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

cnn.add(tf.keras.layers.Conv2D(filters=512, kernel_size=3, padding='same',
    , activation='relu'))
cnn.add(tf.keras.layers.Conv2D(filters=512, kernel_size=3, activation='relu'))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))
```

Add Dropout layer to reduce overfitting

```
cnn.add(tf.keras.layers.Dropout(0.25))
```

Flatten the input

```
cnn.add(tf.keras.layers.Flatten())
```

Add Dense layer with 1500 units and ReLU activation

```
cnn.add(tf.keras.layers.Dense(units=1500, activation='relu'))
```

Add Dropout layer to reduce overfitting

```
cnn.add(tf.keras.layers.Dropout(0.4))
```

Add output Dense layer with 38 units and softmax activation

```
cnn.add(tf.keras.layers.Dense(units=38,activation='softmax'))
```

Compile the model with Adam optimizer, categorical crossentropy loss, and accuracy metric

```
cnn.compile(optimizer=tf.keras.optimizers.Adam(
    learning_rate=0.0001), loss='categorical_crossentropy',
    metrics=['accuracy'])
```

Print the model summary

```
cnn.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	
Param #		
conv2d_2 (Conv2D)	(None, 128, 128, 32)	
896		
conv2d_3 (Conv2D)	(None, 126, 126, 32)	
9,248		
max_pooling2d_1 (MaxPooling2D)	(None, 63, 63, 32)	
0		
conv2d_4 (Conv2D)	(None, 63, 63, 64)	
18,496		
conv2d_5 (Conv2D)	(None, 61, 61, 64)	
36,928		
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 64)	
0		

conv2d_6 (Conv2D)	(None, 30, 30, 128)	
73,856		
conv2d_7 (Conv2D)	(None, 28, 28, 128)	
147,584		
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 128)	
0		
conv2d_8 (Conv2D)	(None, 14, 14, 256)	
295,168		
conv2d_9 (Conv2D)	(None, 12, 12, 256)	
590,080		
max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 256)	
0		
conv2d_10 (Conv2D)	(None, 6, 6, 512)	
1,180,160		
conv2d_11 (Conv2D)	(None, 4, 4, 512)	
2,359,808		
max_pooling2d_5 (MaxPooling2D)	(None, 2, 2, 512)	
0		
dropout (Dropout)	(None, 2, 2, 512)	
0		
flatten (Flatten)	(None, 2048)	
0		
dense (Dense)	(None, 1500)	
3,073,500		

dropout_1 (Dropout)	(None, 1500)	
dense_1 (Dense)	(None, 38)	

Total params: 7,842,762 (29.92 MB)

Trainable params: 7,842,762 (29.92 MB)

Non-trainable params: 0 (0.00 B)

Train the model with the training set and validate with the validation set for 10 epochs

```
training_history =
cnn.fit(x=training_set,validation_data=validation_set,epochs=10)

Epoch 1/10
2197/2197 _____ 1790s 813ms/step - accuracy: 0.3815 -
loss: 2.1848 - val_accuracy: 0.8297 - val_loss: 0.5510
Epoch 2/10
2197/2197 _____ 1721s 783ms/step - accuracy: 0.8337 -
loss: 0.5256 - val_accuracy: 0.8910 - val_loss: 0.3283
Epoch 3/10
2197/2197 _____ 1718s 782ms/step - accuracy: 0.9066 -
loss: 0.2950 - val_accuracy: 0.9081 - val_loss: 0.2818
Epoch 4/10
2197/2197 _____ 1714s 780ms/step - accuracy: 0.9330 -
loss: 0.2066 - val_accuracy: 0.9158 - val_loss: 0.2571
Epoch 5/10
2197/2197 _____ 1706s 776ms/step - accuracy: 0.9523 -
loss: 0.1462 - val_accuracy: 0.9481 - val_loss: 0.1617
Epoch 6/10
2197/2197 _____ 1709s 778ms/step - accuracy: 0.9614 -
loss: 0.1186 - val_accuracy: 0.9544 - val_loss: 0.1408
Epoch 7/10
2197/2197 _____ 1726s 786ms/step - accuracy: 0.9704 -
loss: 0.0927 - val_accuracy: 0.9500 - val_loss: 0.1609
Epoch 8/10
2197/2197 _____ 1731s 788ms/step - accuracy: 0.9726 -
loss: 0.0827 - val_accuracy: 0.9468 - val_loss: 0.1605
Epoch 9/10
2197/2197 _____ 1727s 786ms/step - accuracy: 0.9780 -
loss: 0.0672 - val_accuracy: 0.9644 - val_loss: 0.1213
Epoch 10/10
```

```
2197/2197 ————— 1737s 790ms/step - accuracy: 0.9804 -  
loss: 0.0613 - val_accuracy: 0.9668 - val_loss: 0.1108
```

Training set Accuracy

```
train_loss, train_acc = cnn.evaluate(training_set)  
print('Training accuracy:', train_acc)
```

```
2197/2197 ————— 448s 204ms/step - accuracy: 0.9891 -  
loss: 0.0295  
Training accuracy: 0.9904544949531555
```

Validation set Accuracy

```
val_loss, val_acc = cnn.evaluate(validation_set)  
print('Validation accuracy:', val_acc)
```

```
550/550 ————— 113s 205ms/step - accuracy: 0.9682 -  
loss: 0.1085  
Validation accuracy: 0.9668222069740295
```

Save the trained model to a file

```
cnn.save('trained_plant_disease_model.keras')
```

Return dictionary of training history

```
training_history.history  
{ 'accuracy': [0.5865281820297241,  
0.8562344312667847,  
0.9146454334259033,  
0.9385020136833191,  
0.9554022550582886,  
0.9635109305381775,  
0.9710505604743958,  
0.9749626517295837,  
0.9788320660591125,  
0.9806387424468994],  
'loss': [1.3972781896591187,  
0.45114046335220337,  
0.26727038621902466,  
0.18771201372146606,  
0.13615547120571136,  
0.1123681366443634,  
0.08959349989891052,  
0.07606132328510284,  
0.06665454059839249,
```

```

0.059280991554260254],
'val_accuracy': [0.8296722173690796,
0.8910198211669922,
0.9080924391746521,
0.915831983089447,
0.9480992555618286,
0.954359233379364,
0.9499772191047668,
0.946790337562561,
0.9644320607185364,
0.9668222069740295],
'val_loss': [0.5510382652282715,
0.3282625079154968,
0.2818470895290375,
0.25710904598236084,
0.1616627275943756,
0.1407908797264099,
0.1609172374010086,
0.16052775084972382,
0.12127251923084259,
0.11083154380321503]}

```

Recording History in json

```

import json
with open('training_hist.json', 'w') as f:
    json.dump(training_history.history, f)

```

Print the keys of the training history dictionary

```

print(training_history.history.keys())

dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])

```


Plot training and validation accuracy over epochs

```

num_epochs = len(training_history.history['accuracy'])
epochs = range(1, num_epochs + 1)

# Plot training and validation accuracy
plt.plot(epochs, training_history.history['accuracy'], color='red',
label='Training Accuracy')
plt.plot(epochs, training_history.history['val_accuracy'],
color='blue', label='Validation Accuracy')
plt.xlabel('No. of Epochs')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.show()

```



```
class_name = validation_set.class_names
test_set = tf.keras.utils.image_dataset_from_directory(
    'valid',
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=1,
    image_size=(128, 128),
    shuffle=False,
    seed=None,
    validation_split=None,
    subset=None,
    interpolation="bilinear",
    follow_links=False,
    crop_to_aspect_ratio=False
)
```

Found 17572 files belonging to 38 classes.

```
y_pred = cnn.predict(test_set)
predicted_categories = tf.argmax(y_pred, axis=1)
```


17572/17572 ————— 287s 16ms/step

```
true_categories = tf.concat([y for x, y in test_set], axis=0)
```

```
Y_true = tf.argmax(true_categories, axis=1)
```

Y_true

```
<tf.Tensor: shape=(17572,), dtype=int64, numpy=array([ 0,  0,  0, ..., 37, 37, 37])>
```

predicted_categories

```
<tf.Tensor: shape=(17572,), dtype=int64, numpy=array([ 0,  0,  0, ..., 37, 37, 37])>
```

```
from sklearn.metrics import confusion_matrix, classification_report
cm = confusion_matrix(Y_true, predicted_categories)
```

```
# Precision Recall Fscore
```

```
print(classification_report(Y_true, predicted_categories, target_names=c
lass_name))
```

				precision
recall	f1-score	support		
			Apple___Apple_scab	0.98
0.97	0.97	504		
			Apple___Black_rot	0.97
0.99	0.98	497		
			Apple___Cedar_apple_rust	0.97
0.97	0.97	440		
			Apple___healthy	0.98
0.96	0.97	502		
			Blueberry___healthy	0.93
0.99	0.96	454		
			Cherry_(including_sour)___Powdery_mildew	0.99
0.99	0.99	421		
			Cherry_(including_sour)___healthy	0.98
0.99	0.99	456		
			Corn_(maize)___Cercospora_leaf_spot Gray_leaf_spot	0.86
0.98	0.92	410		
			Corn_(maize)___Common_rust_	0.99
0.99	0.99	477		
			Corn_(maize)___Northern_Leaf_Blight	0.98
0.89	0.93	477		
			Corn_(maize)___healthy	1.00
1.00	1.00	465		
			Grape___Black_rot	0.96
0.99	0.98	472		
			Grape___Esca_(Black_Measles)	1.00
0.96	0.98	480		

1.00	Grape___Leaf_blight_(Isariopsis_Leaf_Spot)	0.99
0.99	430	
	Grape___healthy	1.00
0.99	423	
0.98	Orange___Haunglongbing_(Citrus_greening)	0.99
	503	
0.96	Peach___Bacterial_spot	0.98
0.97	459	
0.99	Peach___healthy	0.99
	432	
1.00	Pepper,_bell___Bacterial_spot	0.91
0.95	478	
0.96	Pepper,_bell___healthy	0.98
0.97	497	
0.97	Potato___Early_blight	0.98
	485	
0.94	Potato___Late_blight	0.97
0.95	485	
0.97	Potato___healthy	0.99
0.98	456	
0.99	Raspberry___healthy	0.98
0.99	445	
0.97	Soybean___healthy	1.00
0.98	505	
1.00	Squash___Powdery_mildew	0.98
0.99	434	
0.98	Strawberry___Leaf_scorch	0.98
0.98	444	
0.99	Strawberry___healthy	0.99
0.99	456	
0.98	Tomato___Bacterial_spot	0.92
0.95	425	
0.92	Tomato___Early_blight	0.89
0.90	480	
0.86	Tomato___Late_blight	0.98
0.91	463	
0.99	Tomato___Leaf_Mold	0.94
0.96	470	
0.89	Tomato___Septoria_leaf_spot	0.90
0.89	436	
0.94	Tomato___Spider_mites Two-spotted_spider_mite	0.94
0.94	435	
0.84	Tomato___Target_Spot	0.97
0.90	457	
1.00	Tomato___Tomato_Yellow_Leaf_Curl_Virus	0.94
0.97	490	
1.00	Tomato___Tomato_mosaic_virus	0.97
0.99	448	
	Tomato___healthy	0.99

0.99	0.99	481		
			accuracy	
0.97	17572			
			macro avg	0.97
0.97	0.97	17572		
			weighted avg	0.97
0.97	0.97	17572		

```
plt.figure(figsize=(40, 40))
sns.heatmap(cm,annot=True,annot_kws={"size": 10})

plt.xlabel('Predicted Class',fontsize = 20)
plt.ylabel('Actual Class',fontsize = 20)
plt.title('Plant Disease Prediction Confusion Matrix',fontsize = 25)
plt.show()
```

Hant Disease Prediction Confusion Matrix

This heatmap displays the confusion matrix for a Hant Disease prediction model. The x-axis represents the Predicted Class (0 to 37) and the y-axis represents the Actual Class (0 to 37). The diagonal elements, which represent correct classifications, are labeled with their respective counts. The color scale on the right indicates the magnitude of the counts, ranging from 0 (dark purple) to 400 (light orange).

Key observations from the matrix:

- The diagonal elements are labeled with counts, such as $4.9e+02$ for class 0, $1.0e+01$ for class 1, and $1.0e+01$ for class 37.
- The matrix shows a strong diagonal, indicating high predictive accuracy across most classes.
- There are some off-diagonal elements, particularly in the lower right quadrant, suggesting some misclassification for higher-indexed classes.